



**KubeCon**



**CloudNativeCon**

**Europe 2023**



KubeCon

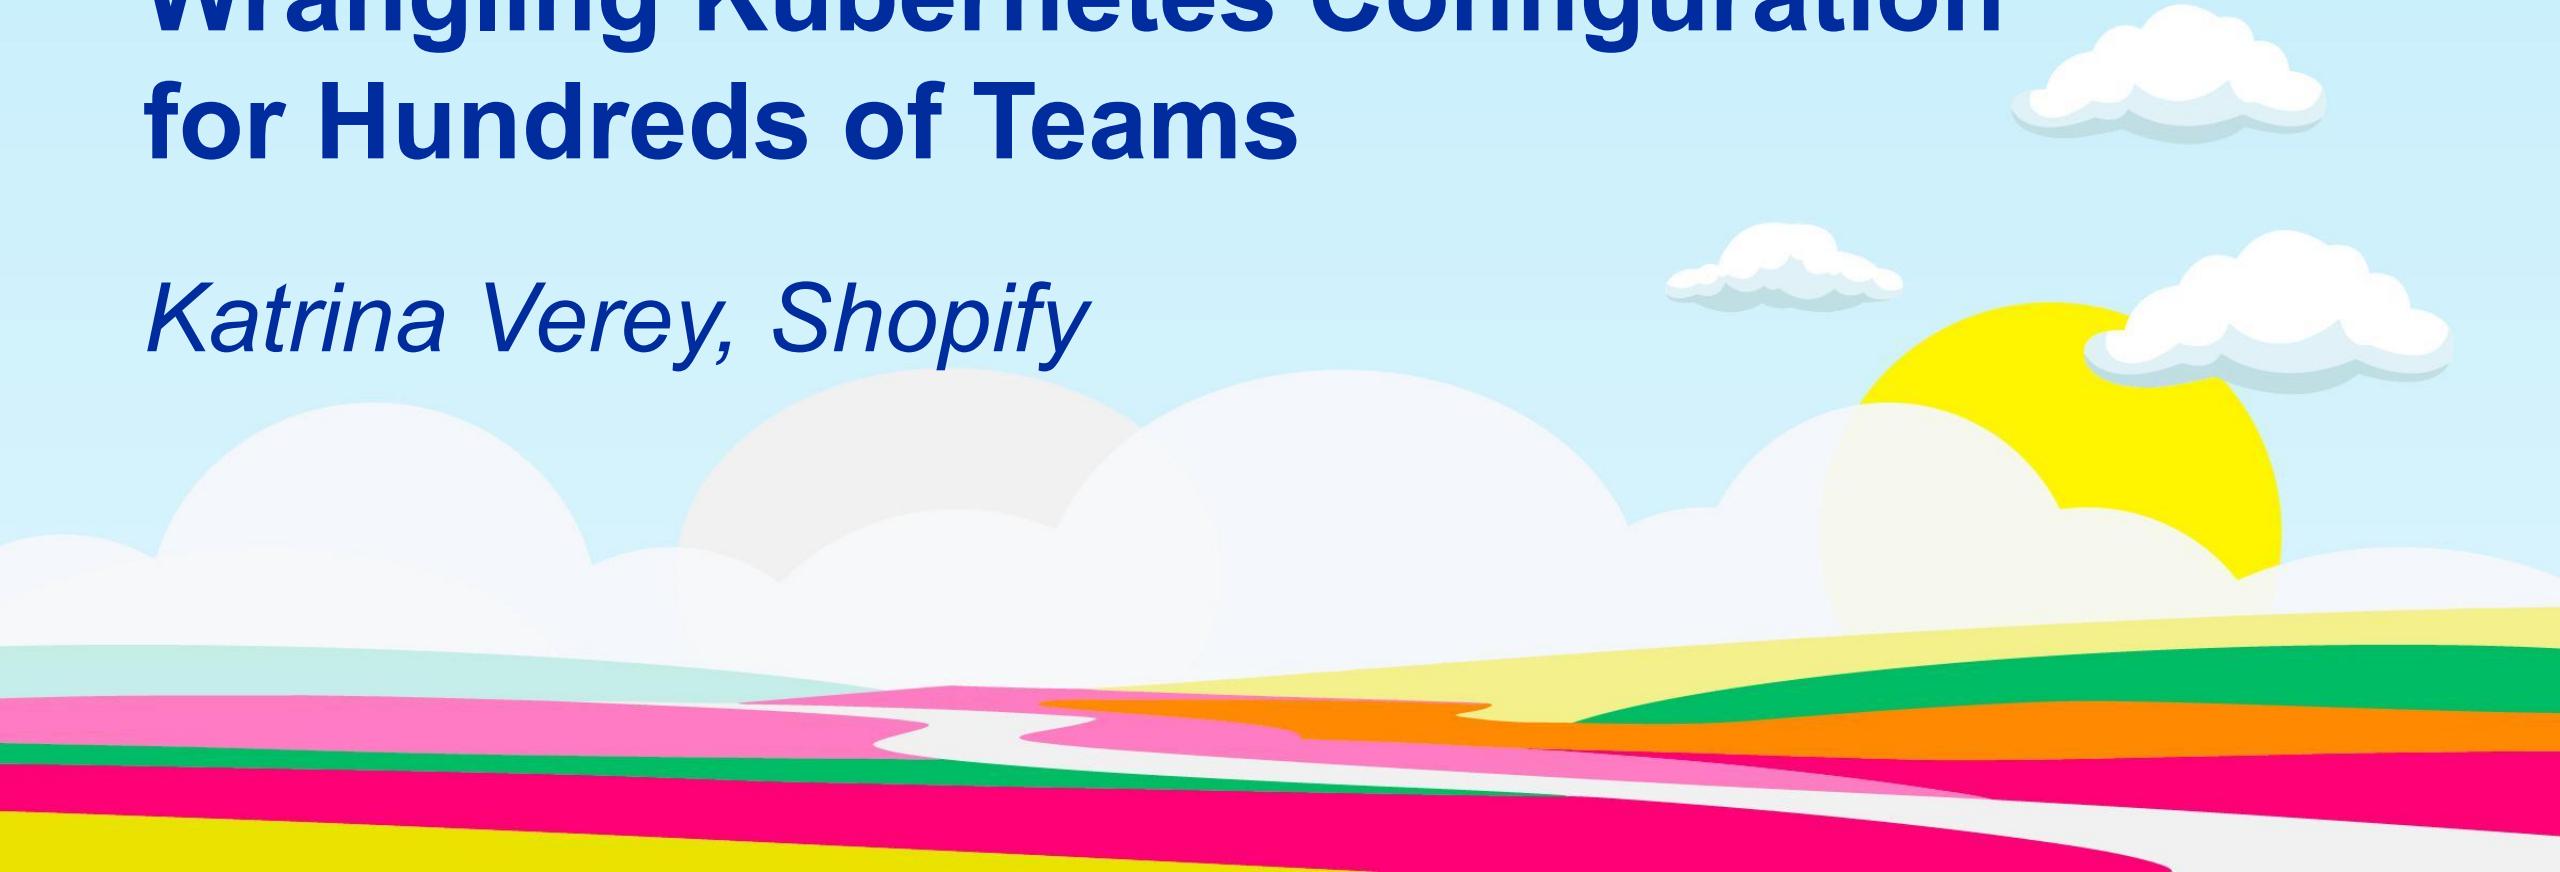


CloudNativeCon

Europe 2023

# 1M Lines of YAML: Wrangling Kubernetes Configuration for Hundreds of Teams

*Katrina Verey, Shopify*



# About me

- Senior Staff Developer, Production Engineering @ Shopify
- Kubernetes end user and enthusiast since 2016
- Upstream lead since 2021:
  - SIG CLI Tech Lead & Chair
  - Kustomize subproject owner
  - KRM Functions subproject owner



# About this talk

- ✓ Managing Kubernetes YAML at scale
  - 1500+ app configuration sets ~> 1M+ lines of YAML
  - Far more than that in our 350+ production clusters*
- ✓ Platform Engineering principles
- ✗ Specific config management tools
- ✓ Systems those tools are part of



# The Problem: ~~YAML~~ Intent Management



"Kubernetes is not just API-driven, but is ***API-centric***."

- Brian Grant, "The Kubernetes Resource Model (KRM)", 2018  
<https://git.k8s.io/design-proposals-archive/architecture/resource-management.md>



## "In Kubernetes, declarative abstractions are primary"

- Brian Grant, "The Kubernetes Resource Model (KRM)", 2018  
(<https://git.k8s.io/design-proposals-archive/architecture/resource-management.md>)

"A simple foundation for a complex space,  
Desired state empowers with effortless grace."

- ChatGPT

# Starting a platform

## Stack

- Rails (default)
- Go

Select the deployments that are needed for your application.

Each selection will be included in a pull request on your service's repo.

You'll have to review this pull request to complete your runtime setup.

For more information visit the [docs](#).

- kube-shell
- elasticsearch
- web
- db-migrate
- asset-uploading
- jobs
- redis
- cloudsql
- memcached

# Per-commit variance

```
spec:  
  containers:  
    - name: web  
      image: "registry.shopify.io/path/to/shopify-app:2f2a598a302060945bda4ffc70b23d6677597936"  
    env:  
      - name: REVISION  
        value: "2f2a598a302060945bda4ffc70b23d6677597936"
```

# Solution: ERB

```
def apply_all_templates
  found = 0
  Dir.foreach(@template_path) do |file|
    file_path = "#{@template_path}/#{file}"
    if File.extname(file) == '.yml'
      found += 1
      run_command('kubectl', 'apply', '-f', file_path, "--namespace=#{@namespace}")
    elsif File.extname(file) == '.erb'
      found += 1
      render_and_apply_template(file_path)
    end
  end
  raise FatalDeploymentError, "No templates found in #{@template_path}" if found.zero?
end
```

# ERB: A big mistake

Hmm, a Turing-complete language  
in config files! WCGW?



*My very smart boss*



It's SO convenient for everyone!  
I'm sure it'll be fine.



*2016 Me*

# Why it was bad

- As predicted, people did all sorts of crazy things with it
  - Pseudo-abstractions? 🤯
  - Complex math? 🤯
  - HTTP calls!?!? 🤯



# Why it was bad

- We couldn't parse it!
  - String scanning as a maintenance strategy 
  - Delayed updates & unintended drift
- Missing interface
  - No separation of concerns
  - Customizations only identifiable via git
  - High learning curve for devs

# Deploy-time rendering: a worse mistake!

Unintended changes can be hard to catch

```
1 + <% %w(web web-perf).each do |name| %>
```

```
1     2     apiVersion: v1
```

```
2     3     kind: Service
```

```
3     4     metadata:
```

```
4       -   name: web
```

```
5     +   name: <%= name %>
```

```
5     6     annotations:
```



```
@@ -231,6 +232,7 @@ spec:
```

```
231 232           - name: star-shopifyio
```

```
232 233           secret:
```

```
233 234           secretName: star-shopifyio
```

```
235 + <% end %>
```

```
234 236     ---
```

# Deploy-time rendering: a worse mistake!

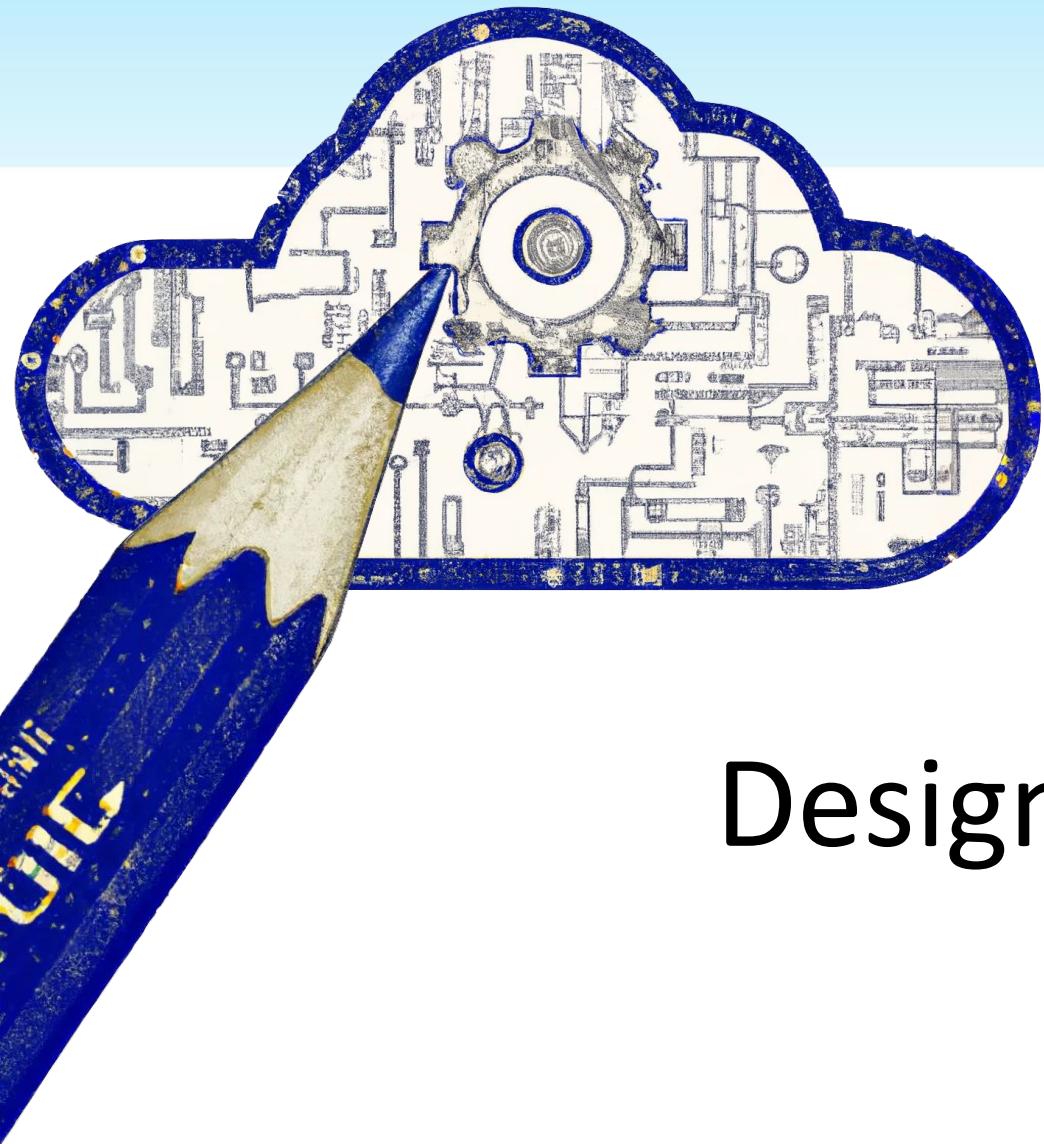
Rollbacks were reliant on rendering

This environment variable was set to (abbreviated):

POD\_IDS: 123456 .....

instead of with commas

POD\_IDS: 1,2,3,4,5,6,7,8,9 .....



# Designing a solution

# Desired properties

## 1. Safety:

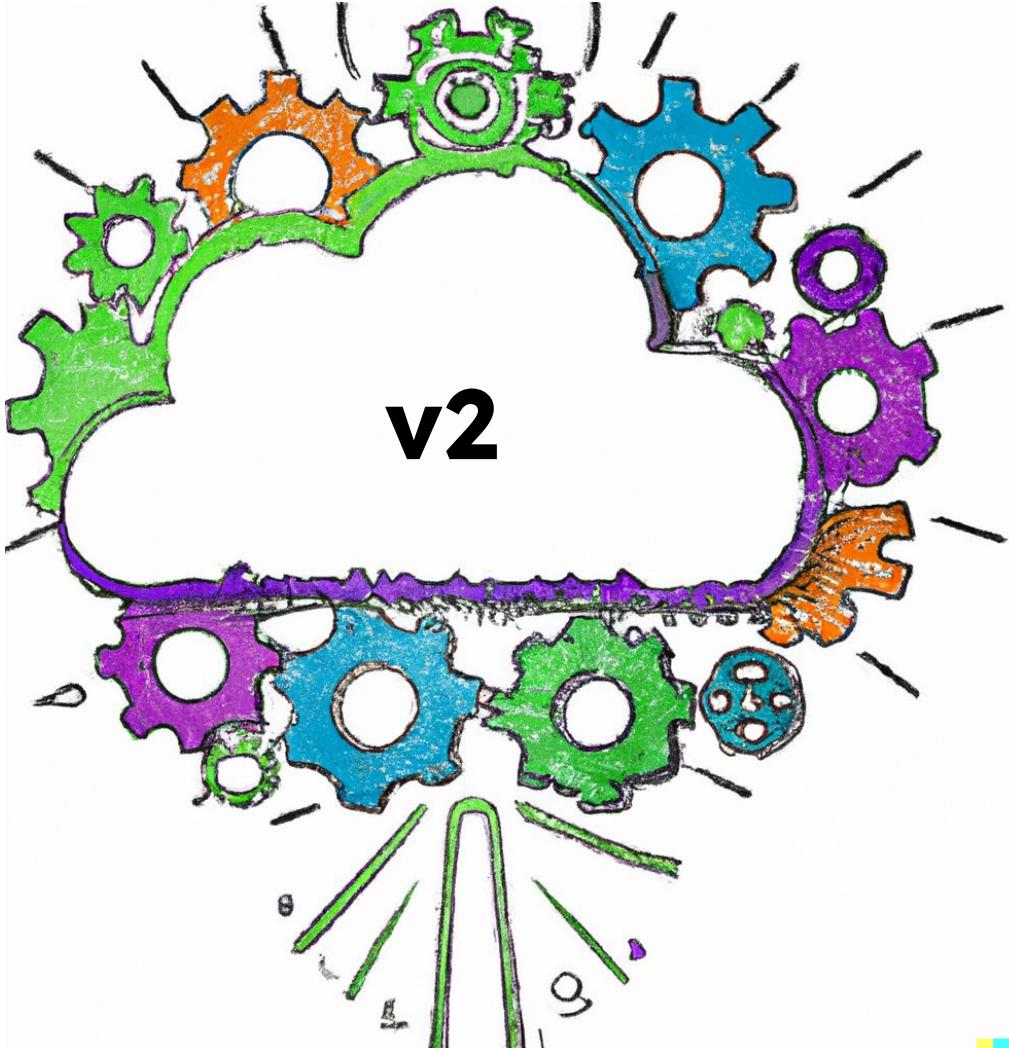
- a. Deploys AND rollbacks must be **repeatable**.
- b. ALL changes must be **easily reviewable**

- 2. Automation-readiness
- 3. Universality



# Desired properties

1. Safety
2. Automation-readiness:
  - a. Possible to **programmatically** make arbitrary updates **reliably**
  - b. Possible to evolve practices through **versioning**
3. Universality



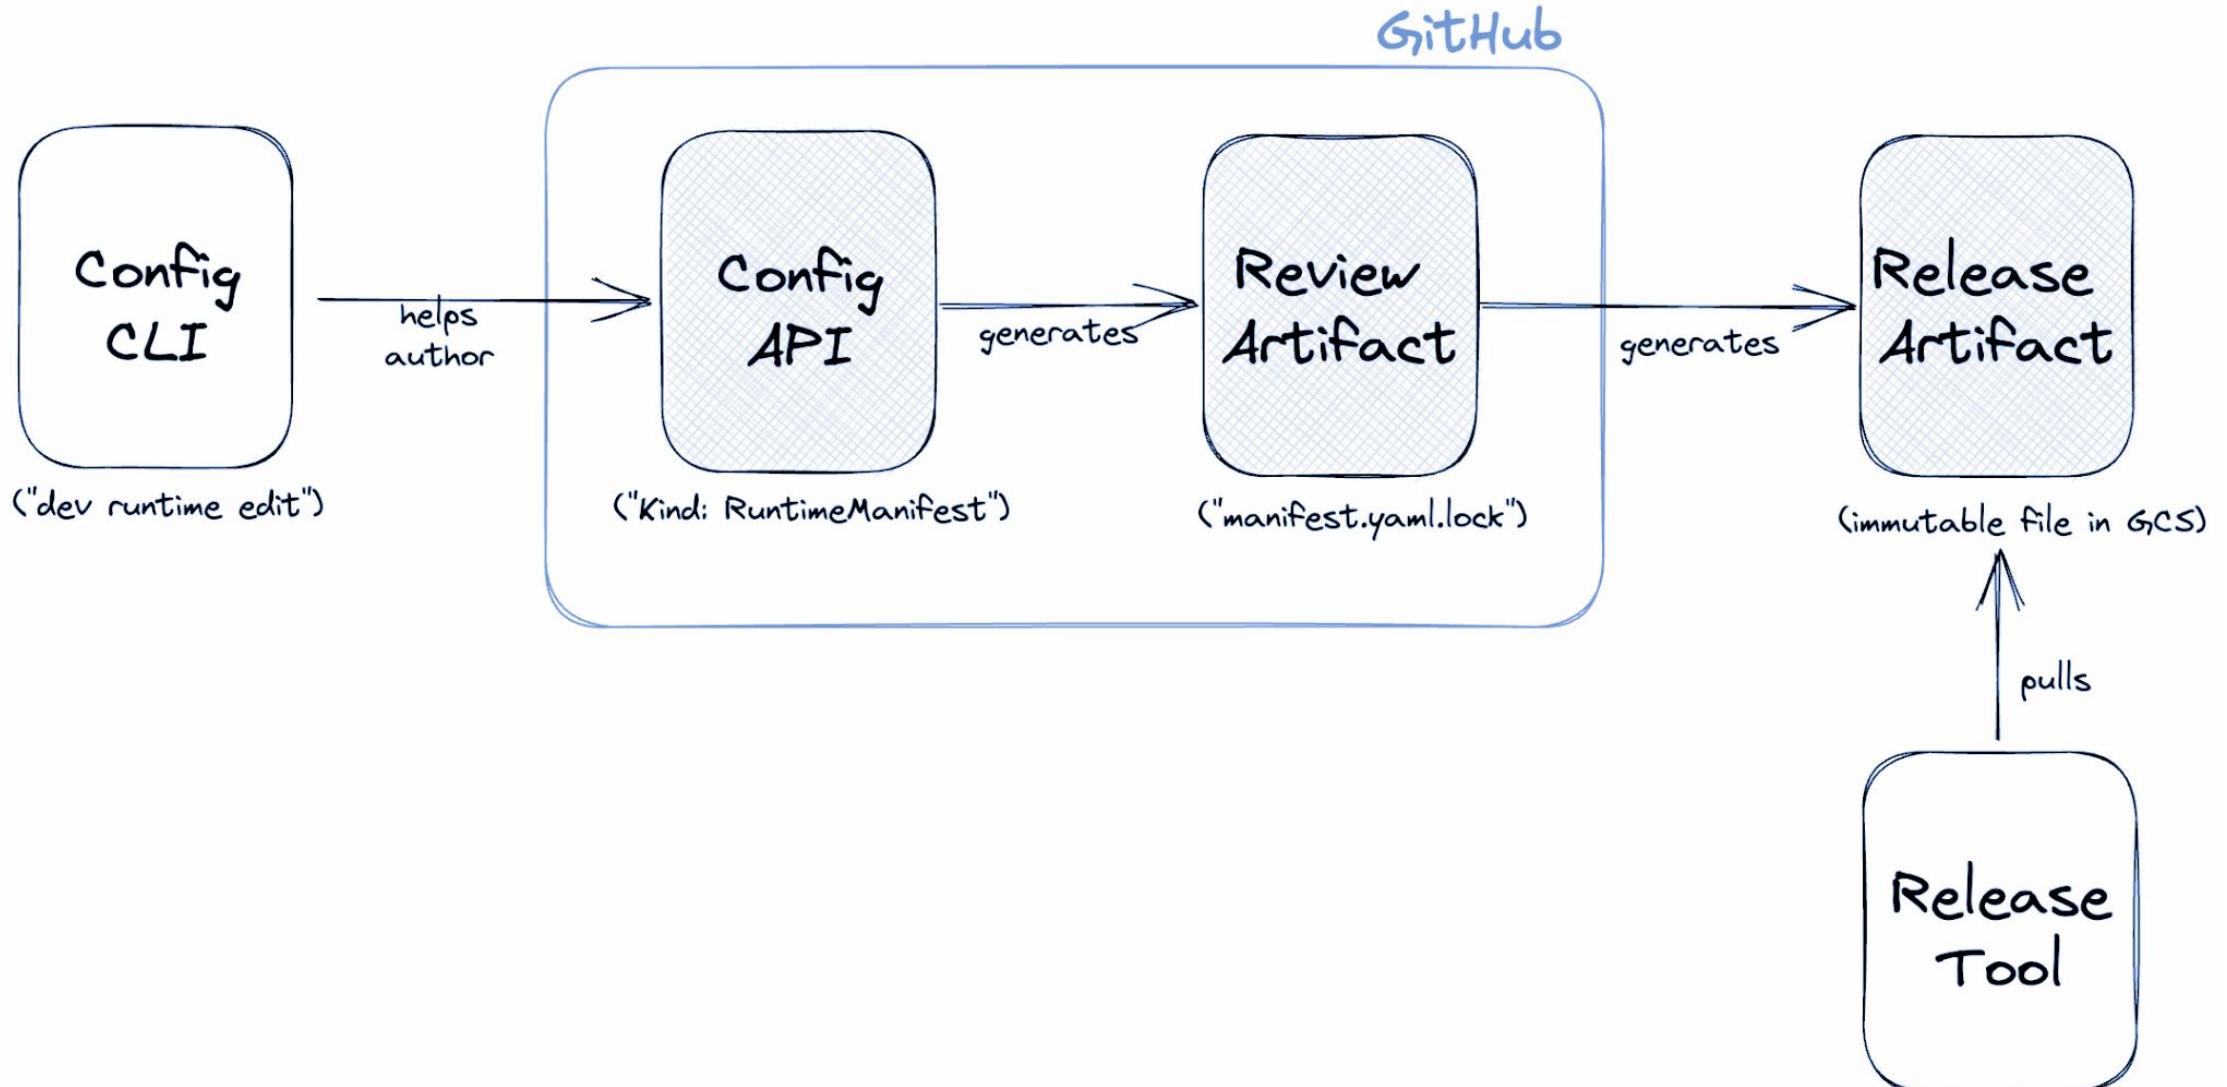
# Desired properties

1. Safety
2. Automation-readiness
3. Universality:
  - a. Simple to start with
  - b. Simple to make common modifications
  - c. Possible to make arbitrary customizations as needed



# What did we build?

# Solution overview



# Config CLI

```
runtime add
```

```
Usage: runtime add [<runtime(s)>] [<component>]
```

Add a component to one or more runtimes.

Multiple runtimes can be manually set using a comma-separated list for the runtime arg

```
runtime edit
```

```
Usage: runtime edit [<runtime>] [<component>]
```

Change the configuration of a component for one or more runtimes

```
runtime upgrade
```

```
Usage: runtime upgrade [<runtime>] [<component>] [--status] [--latest] [--deprecated-versions] [--removed-versions]
```

Upgrade a component for one or more runtimes to the latest version (or specified version).

`all` can be used as an argument for <runtime> or <component> to specify all resource of that type.

# Config API: manifest.yaml

```
1 kind: RuntimeManifest
2 apiVersion: platform.shopify.io/v1alpha1
3 metadata:
4   name: production-unrestricted
5 runtimeInfo:
6   appName: shopify-app
7   env: production
8   appImage: registry.shopify.io/path/to/shopify-app
9 domains:
10  - shopify-app.shopify.io
11 components:
12  - name: web
```

# Config API: manifest.yaml

```
11  components:
12    - name: web
13      basePath: web/v32
14    resources:
15      - base: ingress.yaml
16      - base: certificate.yaml
17      |   name: shopify-app-shopify-io
18      - base: service.yaml
19    patches:
20      - github:Shopify/base_config_repo/patches/legacy/nginx/service.yaml
21    - name: job-worker
22    resources:
23      - base: jobs-deployment.yaml
24      |   name: jobs
25    patches:
26      - repo_local/path/to/custom-labels.yaml
27      - repo_local/path/to/rollout-strategy.yaml
```

# Review artifact: manifest.yaml.lock

```
1 # THIS FILE IS AUTO-GENERATED. DO NOT MODIFY IT MANUALLY.  
2 apiVersion: apps/v1  
3 kind: Deployment  
4 metadata:  
5   name: jobs  
6   labels:  
7     name: jobs  
8     env: production  
9  
10  
11  
12  
13  
14  
15  
16  
17  
18  
19  
20  
21  
22
```

# Review artifact: manifest.yaml.lock

97 - basePath: web/v14

98 + basePath: web/v22



847 - activeDeadlineSeconds: 900

848 automountServiceAccountToken: false

849 containers:

850 - name: command-runner

@@ -857,17 +857,30 @@ spec:

857 - configMapRef:

858 name: application-state

859 env:

847 + activeDeadlineSeconds: 300

848 automountServiceAccountToken: false

849 containers:

850 - name: command-runner

857 - configMapRef:

858 name: application-state

859 env:

860 + - name: STATSD\_USE\_NEW\_CLIENT

861 + value: "1"

862 + - name: KUBE\_POD\_NAME

863 + valueFrom:

864 + fieldRef:

865 + fieldPath: metadata.name

# Release artifact

```
1  # THIS FILE IS AUTO-GENERATED. DO NOT MODIFY IT MANUALLY.  
2  apiVersion: apps/v1  
3  kind: Deployment  
4  metadata:  
5    name: jobs  
6    labels:  
7      name: jobs  
8      env: production  
9  
10  
11  
12  
13  
14  
15  
16  
17  
18  
19  
20  
21  
22
```

# Release artifact

## Config interface

```
5  ↘ runtimeInfo:  
6    appName: shopify-app  
7    env: production  
8    appImage: registry.shopify.io/path/to/shopify-app
```

## Review artifact

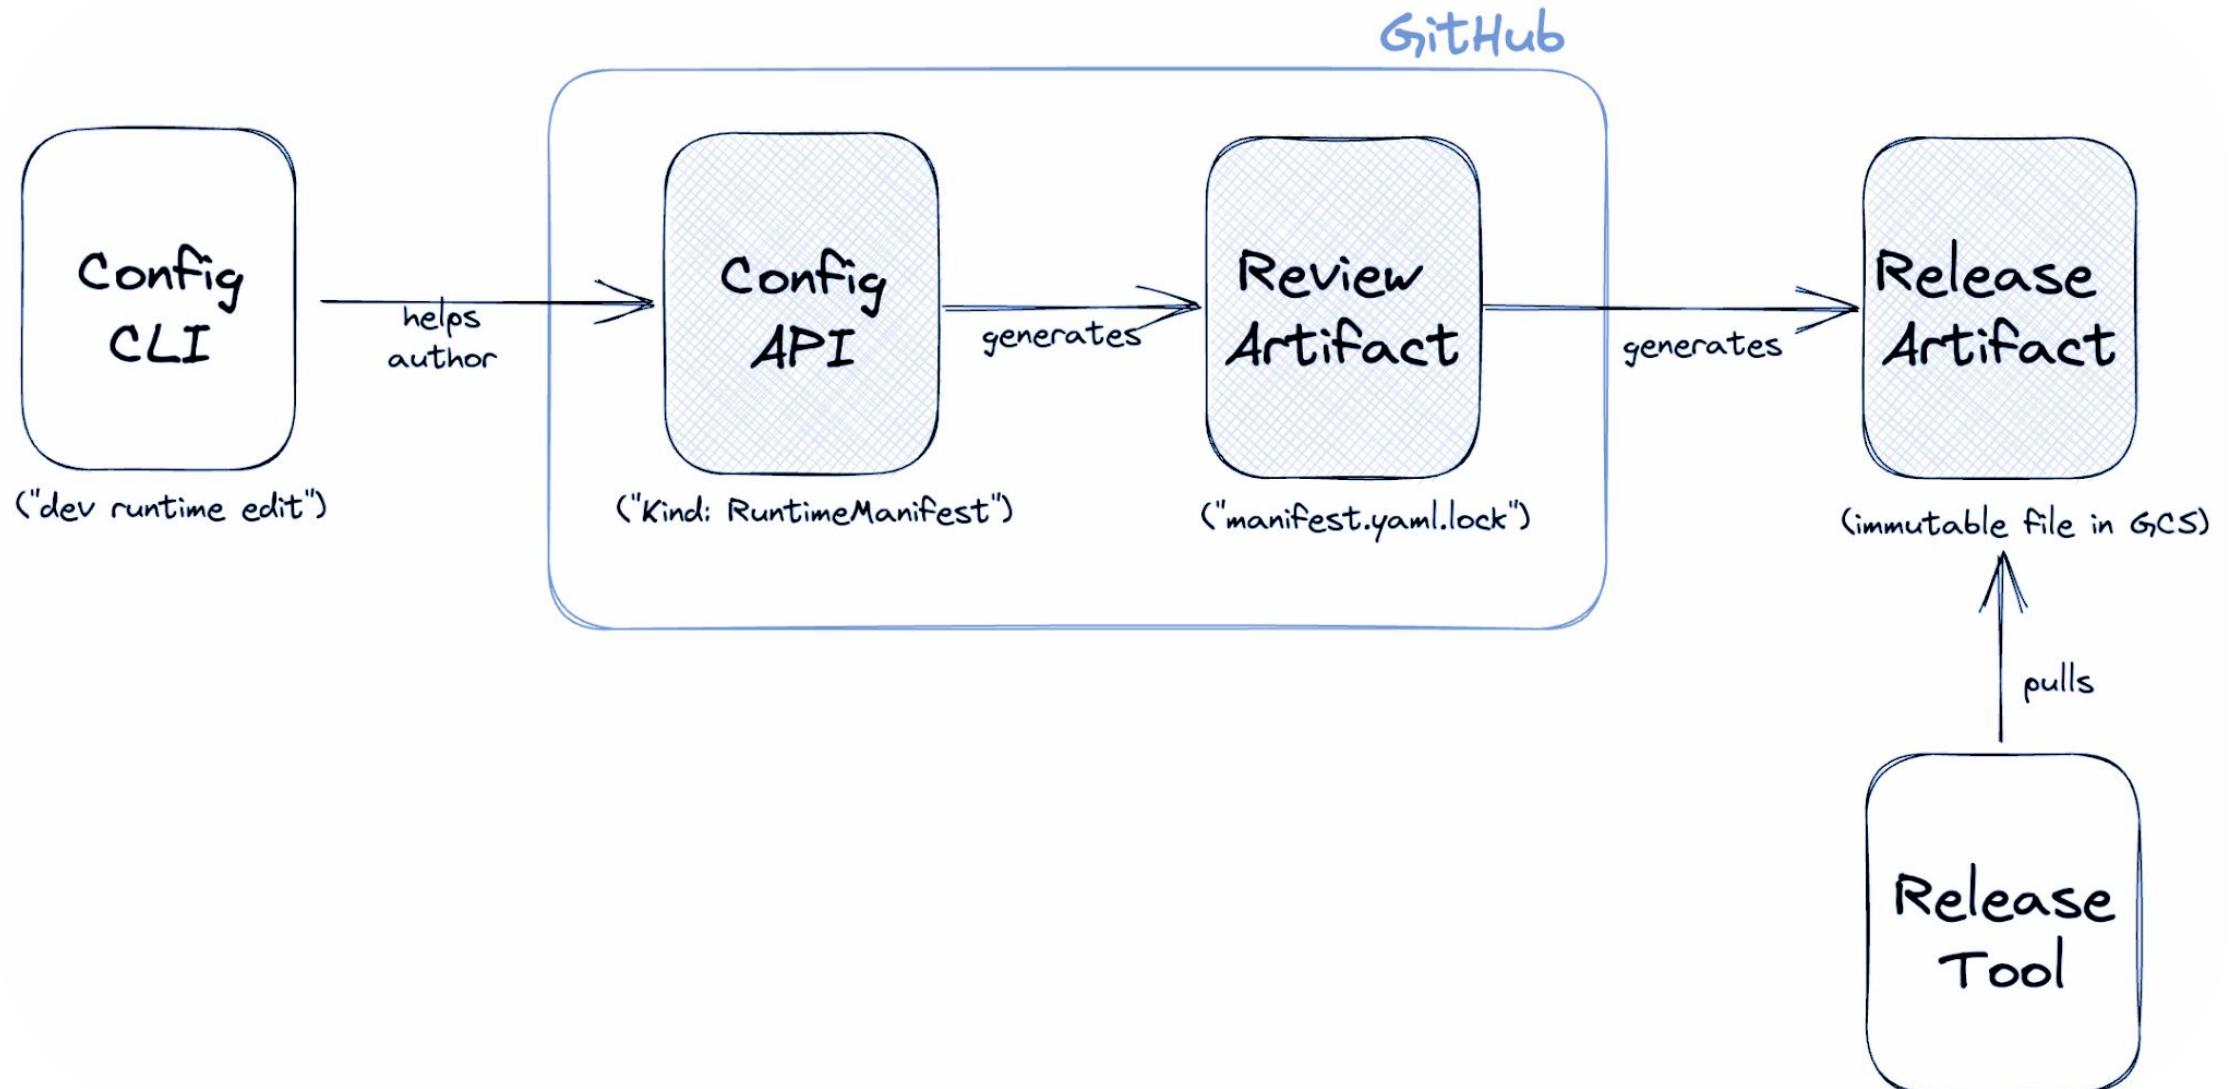
```
↙ metadata:  
↙ annotations:  
    platform.shopify.io/app-containers: jobs
```

## Release artifact

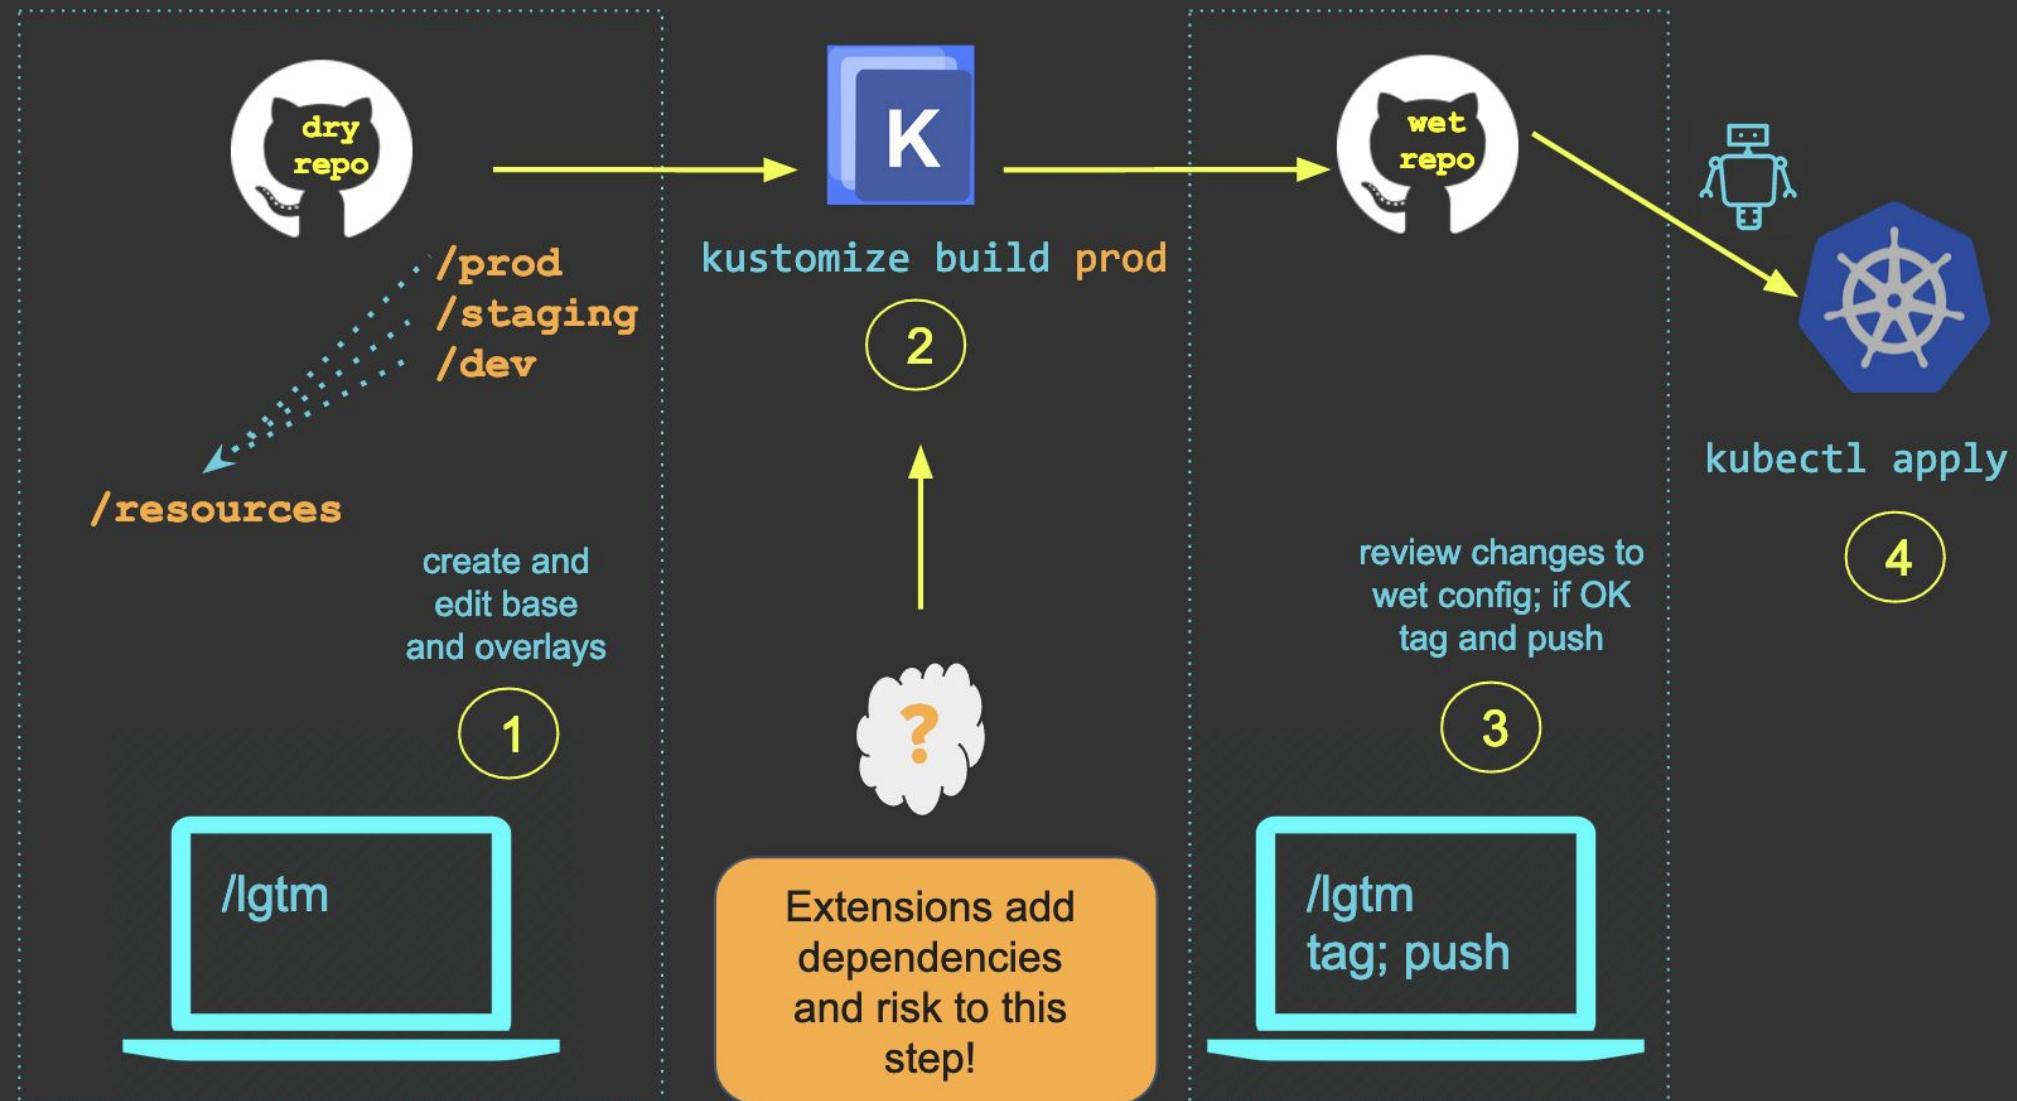
```
spec:  
  containers:  
  - name: jobs  
    image: registry.shopify.io/path/to/shopify-app:2f2a598a302060945bda4ffc70b23d6677597936
```

# What else could this look like?

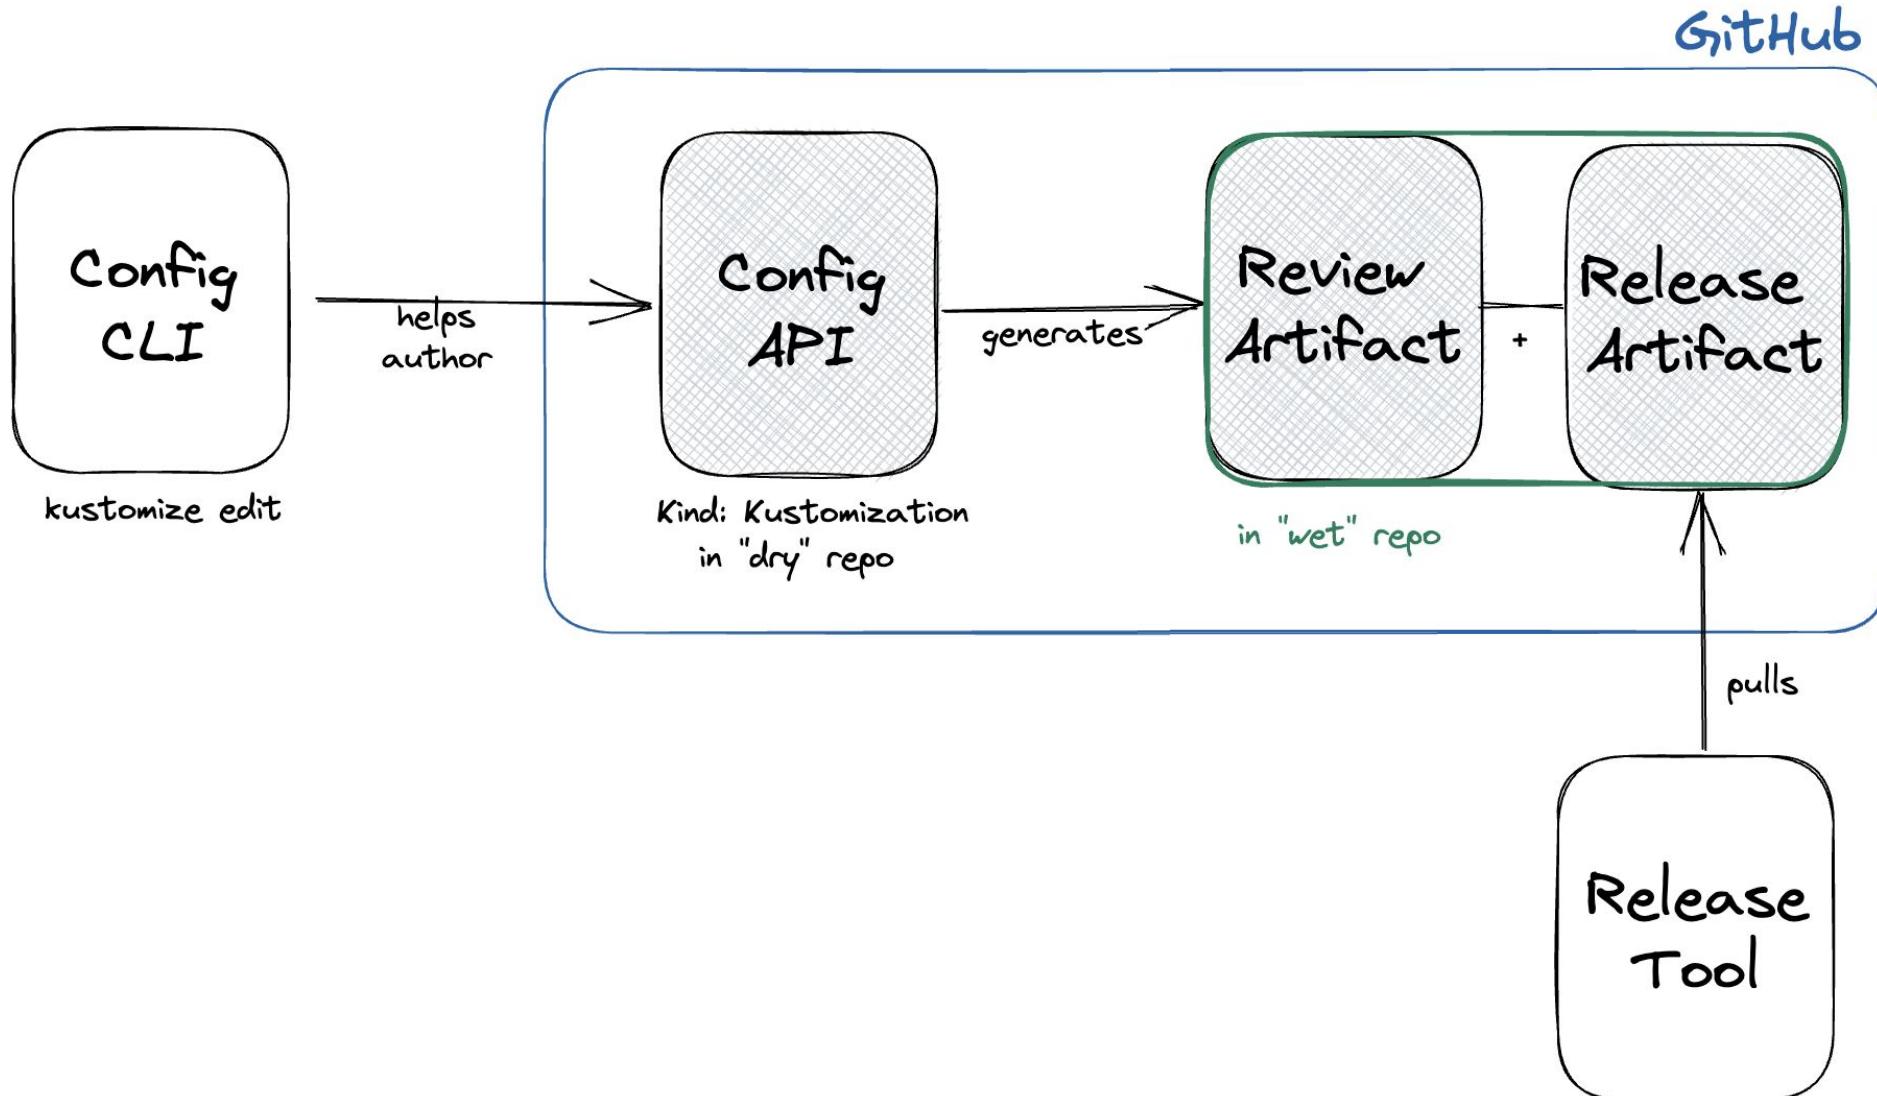
# Alternatives



# Example: Kustomize w/GitOps



# Example: Kustomize w/GitOps



# Why does this work well?

# Principle 1:

## Give your users a way to express *their* intent

# Principle 2:

## If you can't round-trip it, don't distribute it

# Principle 3:

## Make consequential changes visible

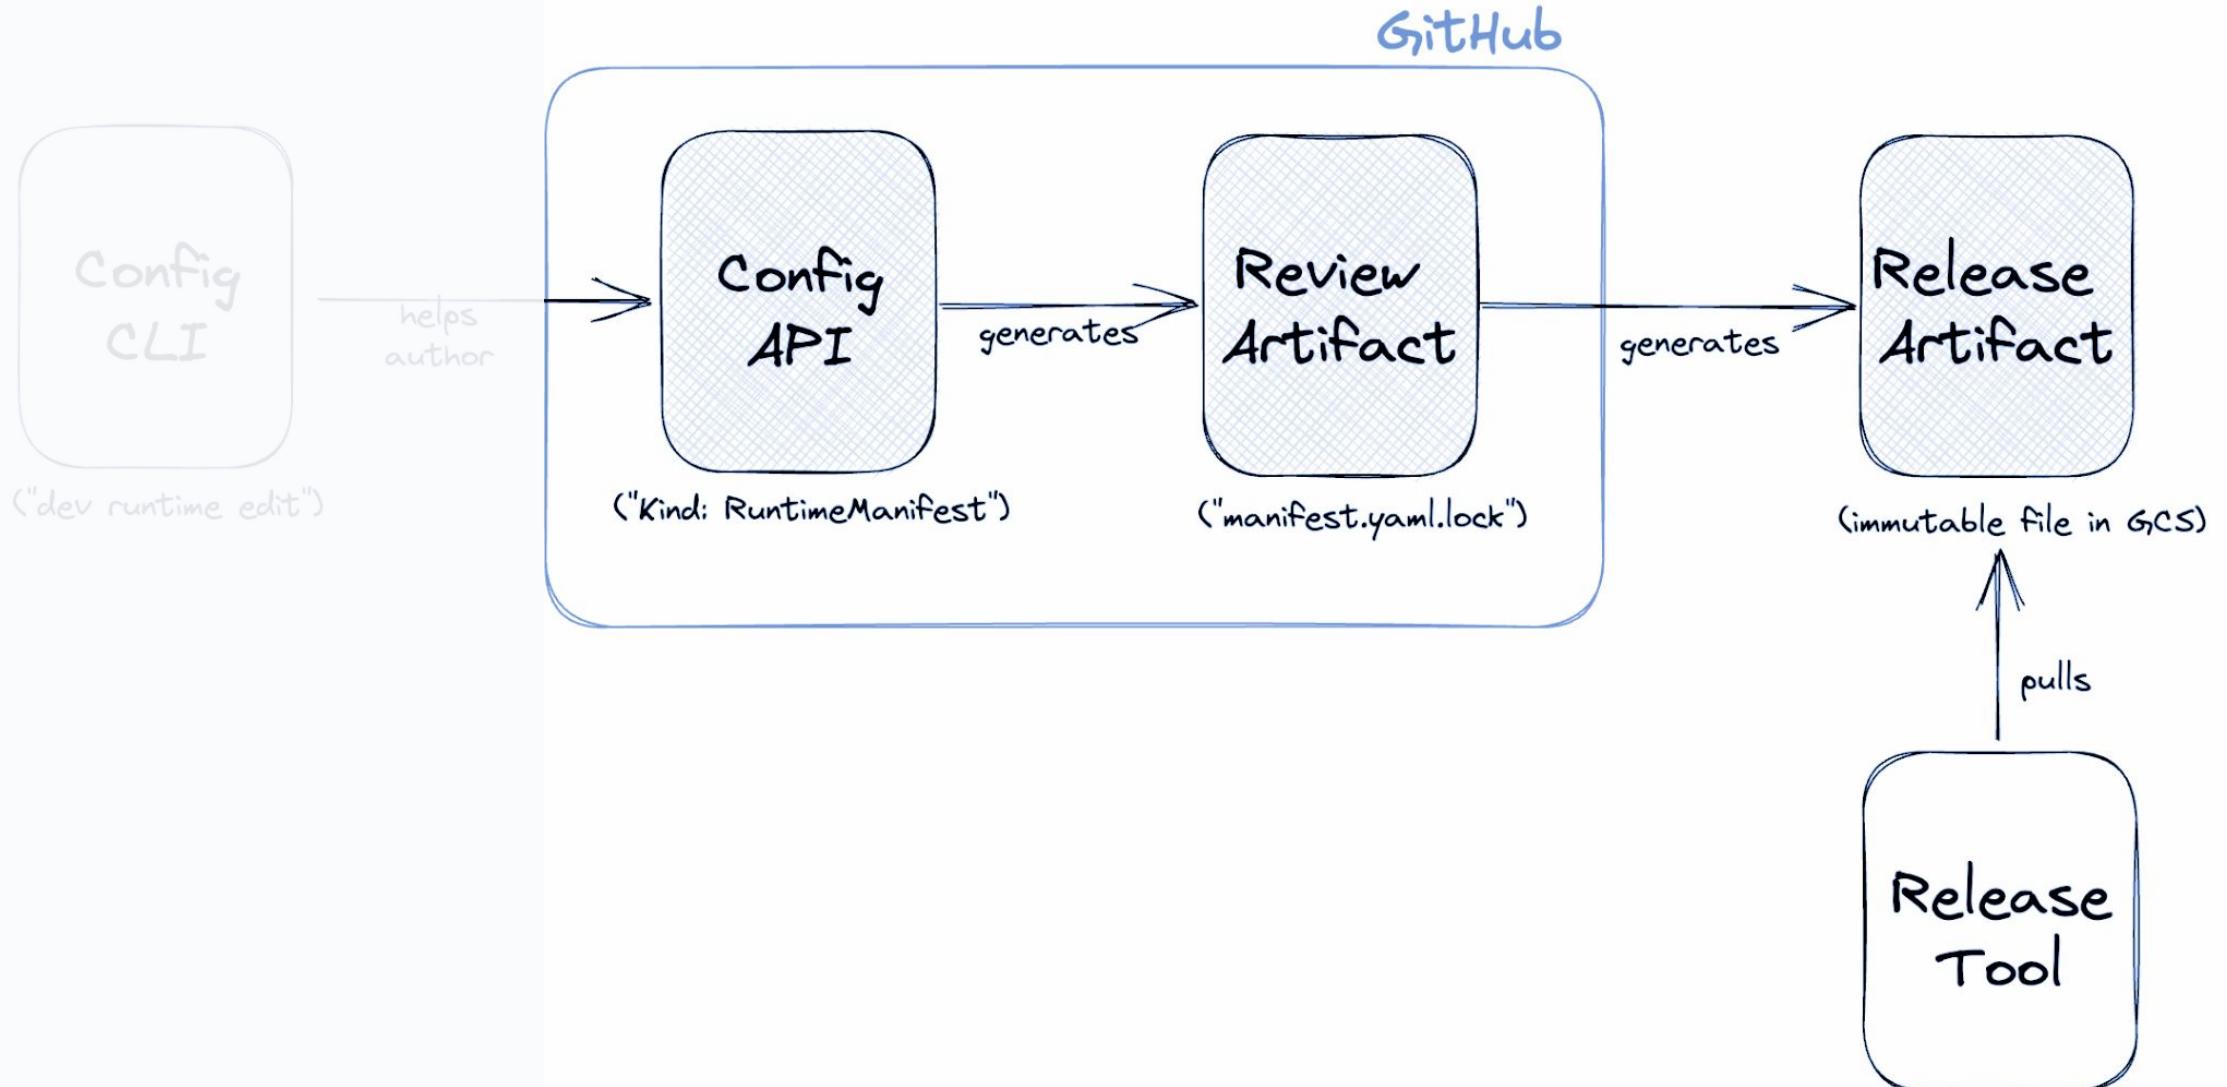
# Principle 4:

## Create release artifacts as early as possible

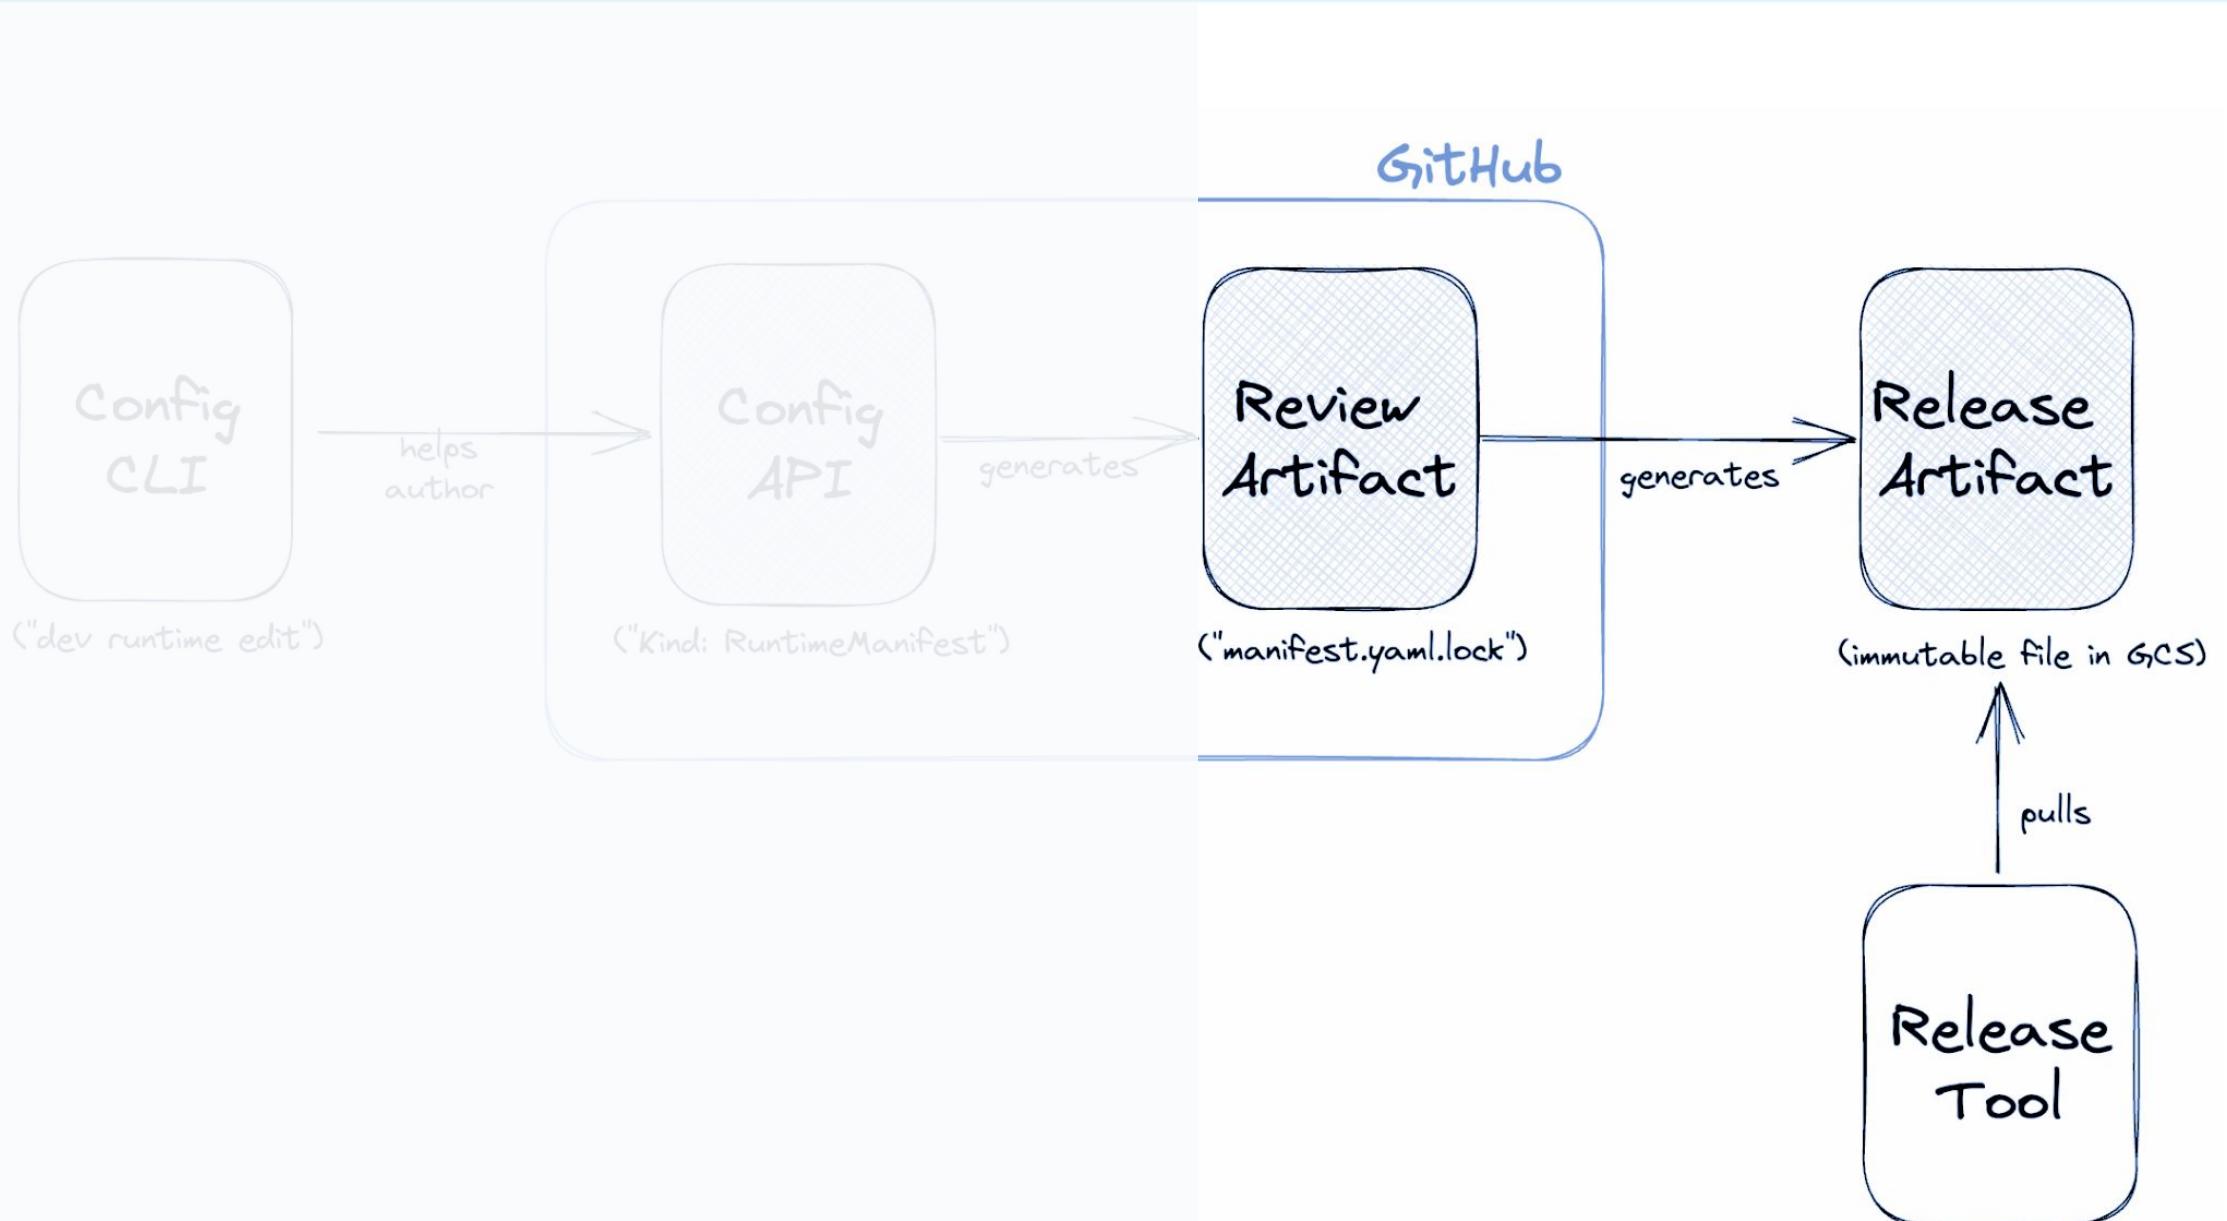
# Principle 5:

## Make it modular

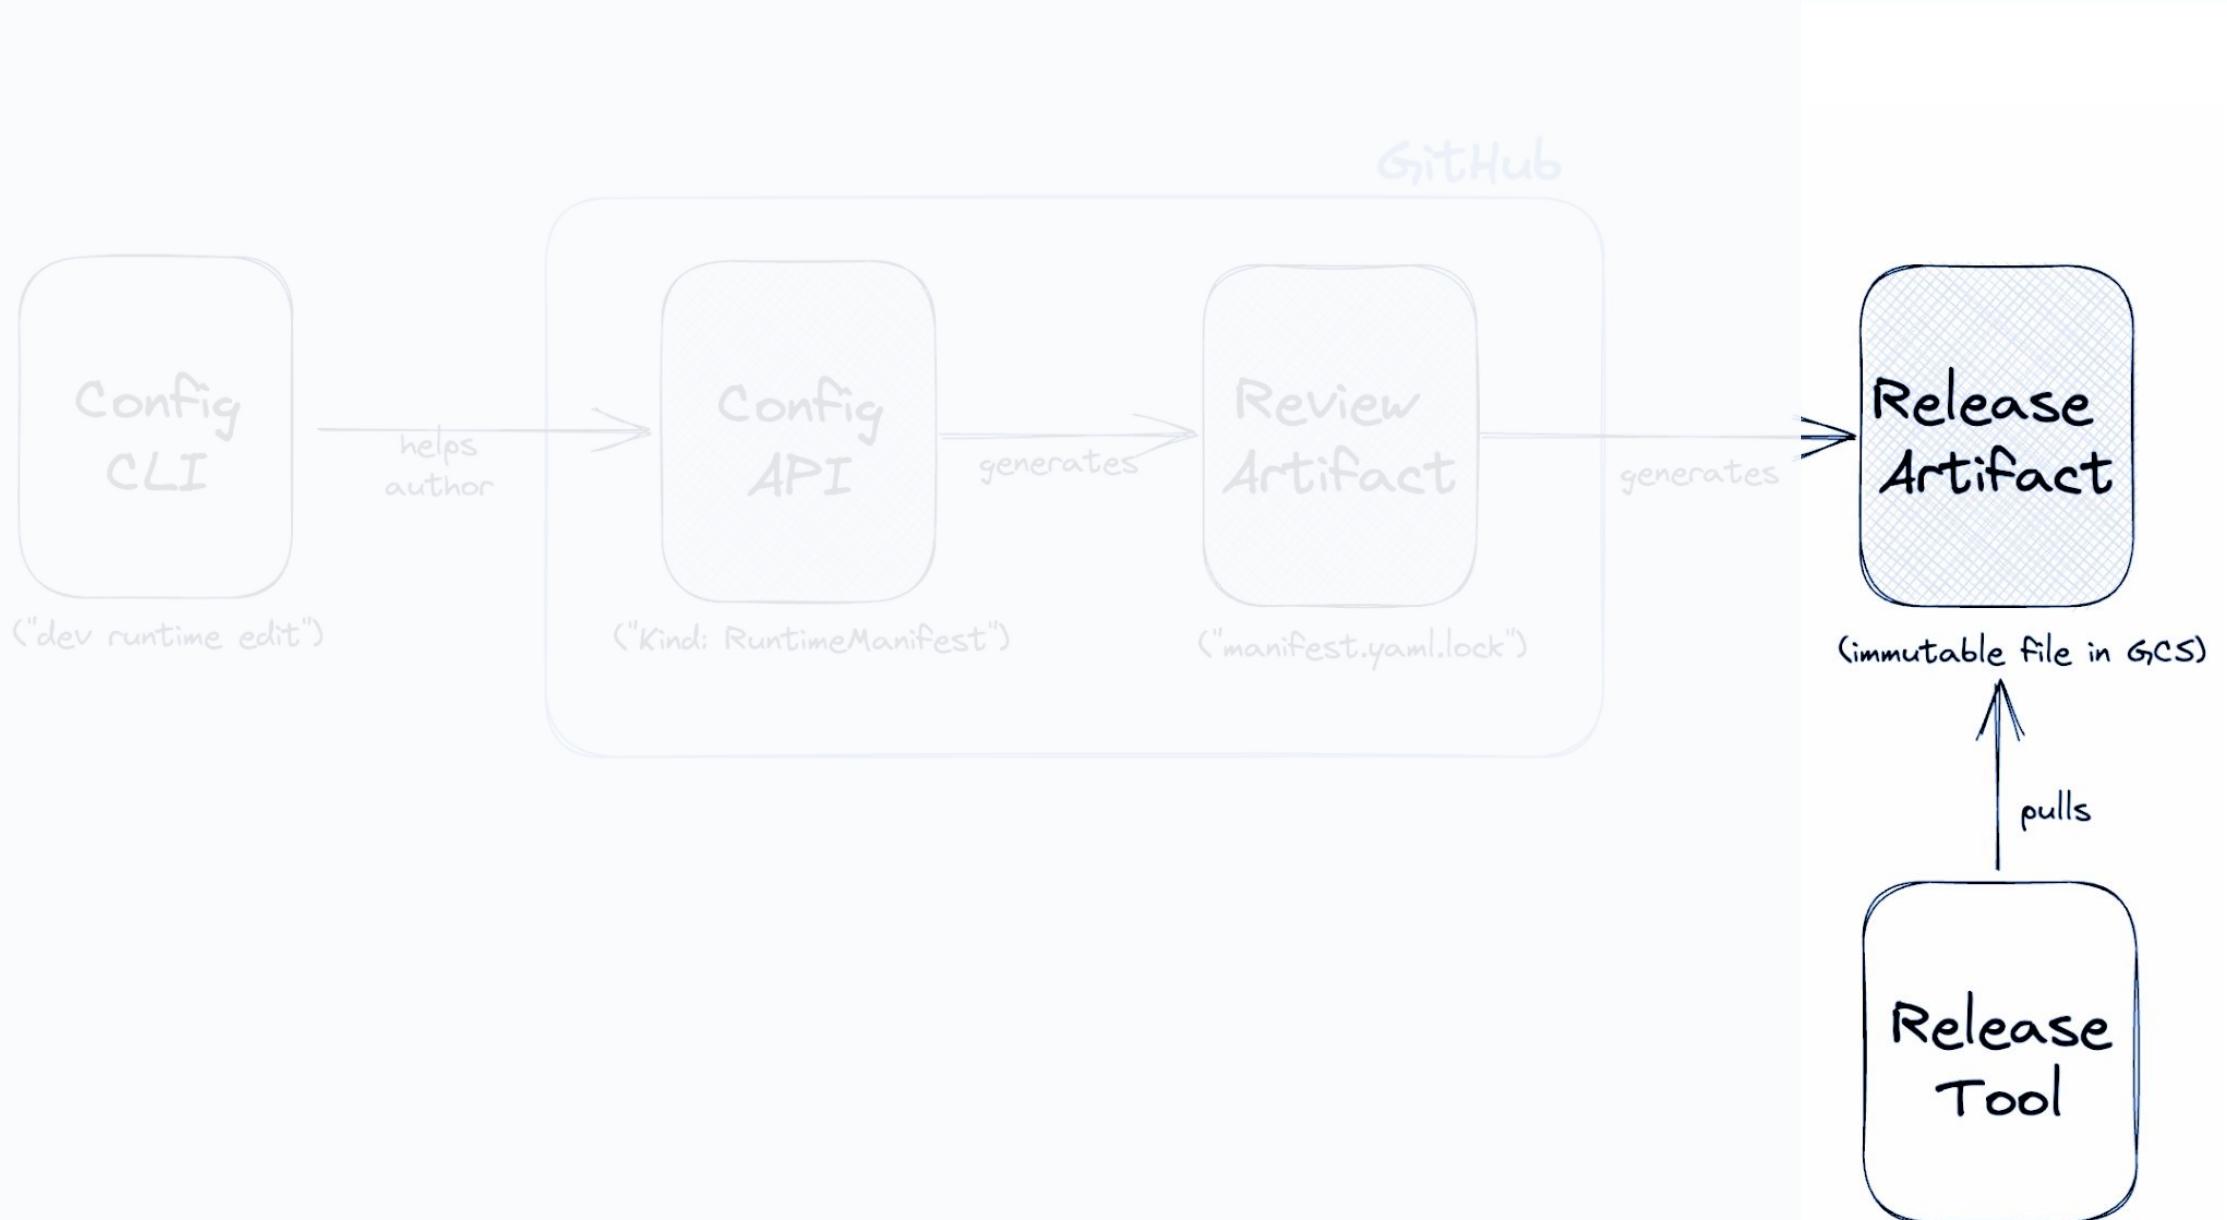
# Modularity

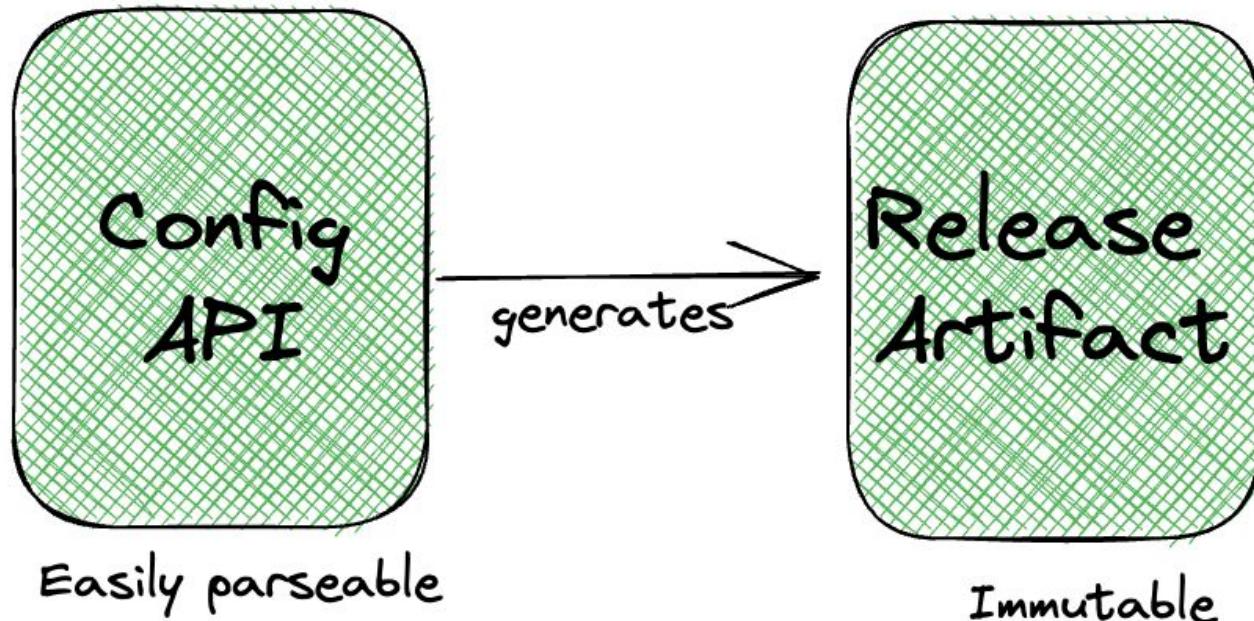


# Modularity



# Modularity





# KRM Functions Framework: A Config API Toolkit

# KRM what?

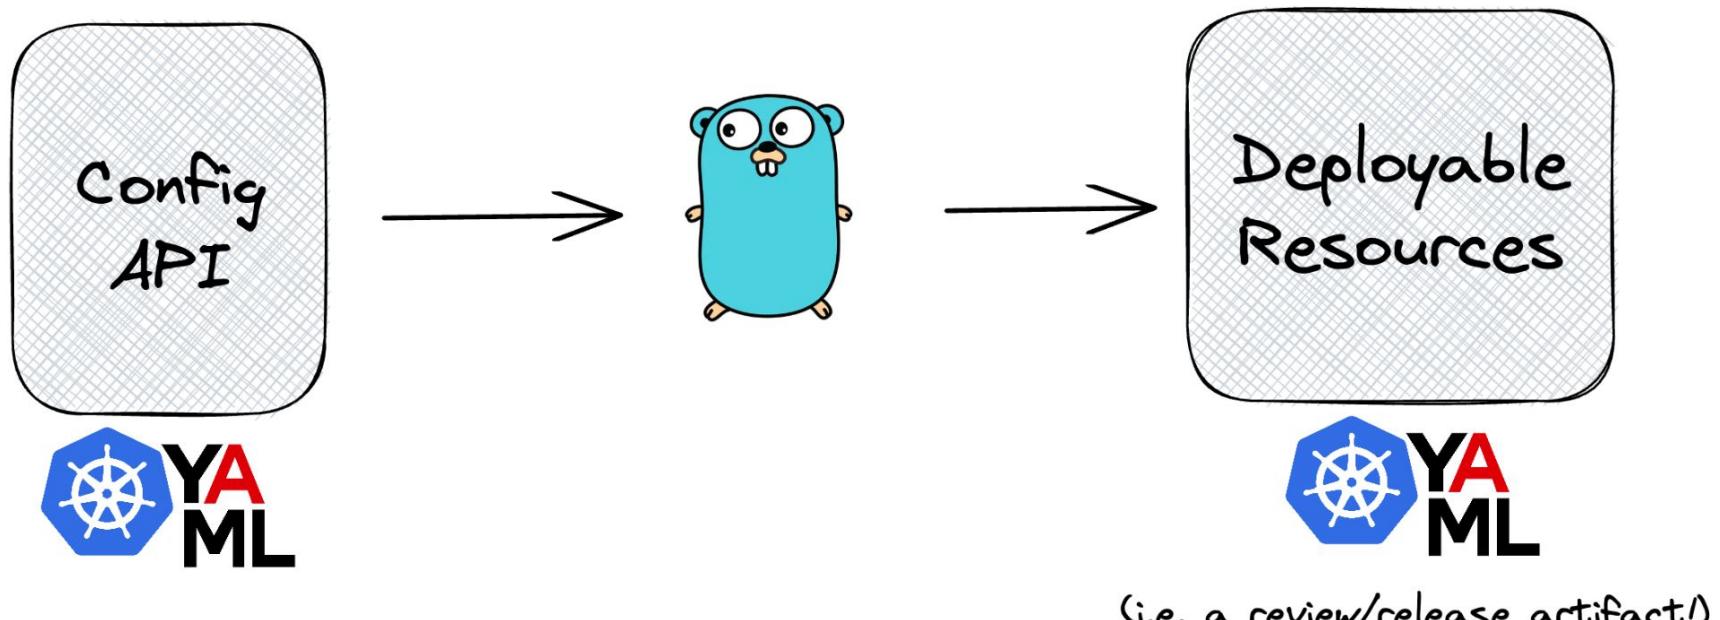
- **KRM:** Kubernetes Resource Model

```
apiVersion: example.com/v1
kind: ExampleApp
metadata:
  name: my-app
```

- **KRM Function:** a piece of code that performs CRUD operations on *local* Kubernetes-style resources based on the desired state declared in a KRM resource.

- **KRM:** Kubernetes Resource Model
- **KRM Function:** a piece of code that performs CRUD operations on *local* Kubernetes-style resources based on the desired state declared in a local Kubernetes-style config API.
- **KRM Functions Specification:** a standard for the design of KRM functions that are small, interoperable, language-independent and chainable as *part of a pipeline that generates fully rendered configurations* that can be applied to a control plane.
- **KRM Functions Framework:** a toolkit for building specification-compliant KRM functions in Go

**KRM Functions Framework:** a toolkit for writing Go code that generates and/or modifies local Kubernetes configuration based on the desired state declared in a local Kubernetes-style API



# KRM Functions Framework

```
1  apiVersion: platform.example.com/v1alpha1
2  kind: ExampleApp
3  metadata:
4      name: simple-app-sample
5  env: production
6  workloads:
7      webWorkers:
8          - name: web-worker
9          domains:
10             - example.com
```

# KRM Functions Framework

```
26 ◻   type ExampleApp struct {
27     // Embedding these structs is required to use controller-gen to produce the CRD
28     metav1.TypeMeta `json:",inline"`
29     metav1.ObjectMeta `json:"metadata"`
30
31     // +kubebuilder:validation:Enum=production;staging;development
32     Env string `json:"env" yaml:"env"`
33
34     // +optional
35     AppImage string `json:"appImage" yaml:"appImage"`
36
37     Workloads Workloads `json:"workloads" yaml:"workloads"`
38
39     // +optional
40     Datastores Datastores `json:"datastores,omitempty" yaml:"datastores,omitempty"`
41
42     // +optional
43     Overrides Overrides `json:"overrides,omitempty" yaml:"overrides,omitempty"`
44 }
```

# KRM Functions Framework

```
21 func (a *ExampleApp) Schema() (*spec.Schema, error) {
22     schema, err := framework.SchemaFromFunctionDefinition(resid.NewGvk(Group, Version, Kind), CRDString)
23     return schema, errors.WrapPrefixf(err, "parsing %s schema", Kind)
24 }
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49 ▼ func (a *ExampleApp) Validate() error {
50     seenDomains := make(map[string]bool)
51     for i, worker := range a.Workloads.WebWorkers {
52         for _, domain := range worker.Domains {
53             if seenDomains[domain] {
54                 return errors.Errorf("duplicate domain %q in worker %d", domain, i)
55             }
56             seenDomains[domain] = true
57         }
58     }
59     return nil
60 }
```

```
26 ▼ func (a *ExampleApp) Default() error {
27     if a.AppImage == "" {
28         a.AppImage = fmt.Sprintf("registry.example.com/path/to/%s", a.ObjectMeta.Name)
29     }
```

# KRM Functions Framework

```
62 ✓ func (a ExampleApp) Filter(items []*yaml.RNode) ([]*yaml.RNode, error) {
63     templates := make([]framework.ResourceTemplate, 0)
64     for _, worker := range a.Workloads.JobWorkers {
65         templates = append(templates, framework.ResourceTemplate{
66             Templates:    parser.TemplateFiles("templates/job_worker.template.yaml").FromFS(templateFS),
67             TemplateData: a.jobWorkerTemplateData(worker),
68         })
69     }
70
71     return templates, nil
72 }
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114     items, err := framework.TemplateProcessor{
115         ResourceTemplates: templates,
116         PatchTemplates:   patches,
117     }.Filter(items)
118     if err != nil {
119         return nil, errors.WrapPrefixf(err, "processing templates")
120     }
121 }
```

# KRM Functions Framework

```
21 ✓ func processKnownAPIGroups(rl *framework.ResourceList) error {
22     p := framework.VersionedAPIProcessor{FilterProvider: framework.GVKFilterMap{
23         "ExampleApp": map[string]kio.Filter{
24             "platform.example.com/v1alpha1": &v1alpha1.ExampleApp{},
25         },
26     }}
27     if err := p.Process(rl); err != nil {
28         return errors.Wrap(err)
29     }
30     var err error
31     // FormatFilter sorts the fields in a deterministic order, which makes the output
32     // more suitable for review in version control.
33     rl.Items, err = filters.FormatFilter{UseSchema: true}.Filter(rl.Items)
34     if err != nil {
35         return errors.WrapPrefixf(err, "formatting output")
36     }
37     return nil
38 }
```

# KRM Functions Framework

app-fn path/to/example\_app.yaml

# Q&A time!

## Session resources

<https://bit.ly/kubecon-1M-lines-of-yaml>



## Session feedback

<https://sched.co/1Hycs>



KubeCon



CloudNativeCon

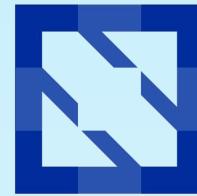
Europe 2023



# 1M Lines of YAML: *Wrangling Kubernetes Configuration for Hundreds of Teams*



KubeCon



CloudNativeCon

Europe 2023

18-21 April



Katrina Verey

Sr. Staff Software Developer

*Shopify*